CircleNet point-attention architecture

Current trained branch · B = batch size · one target circle per l0 patch

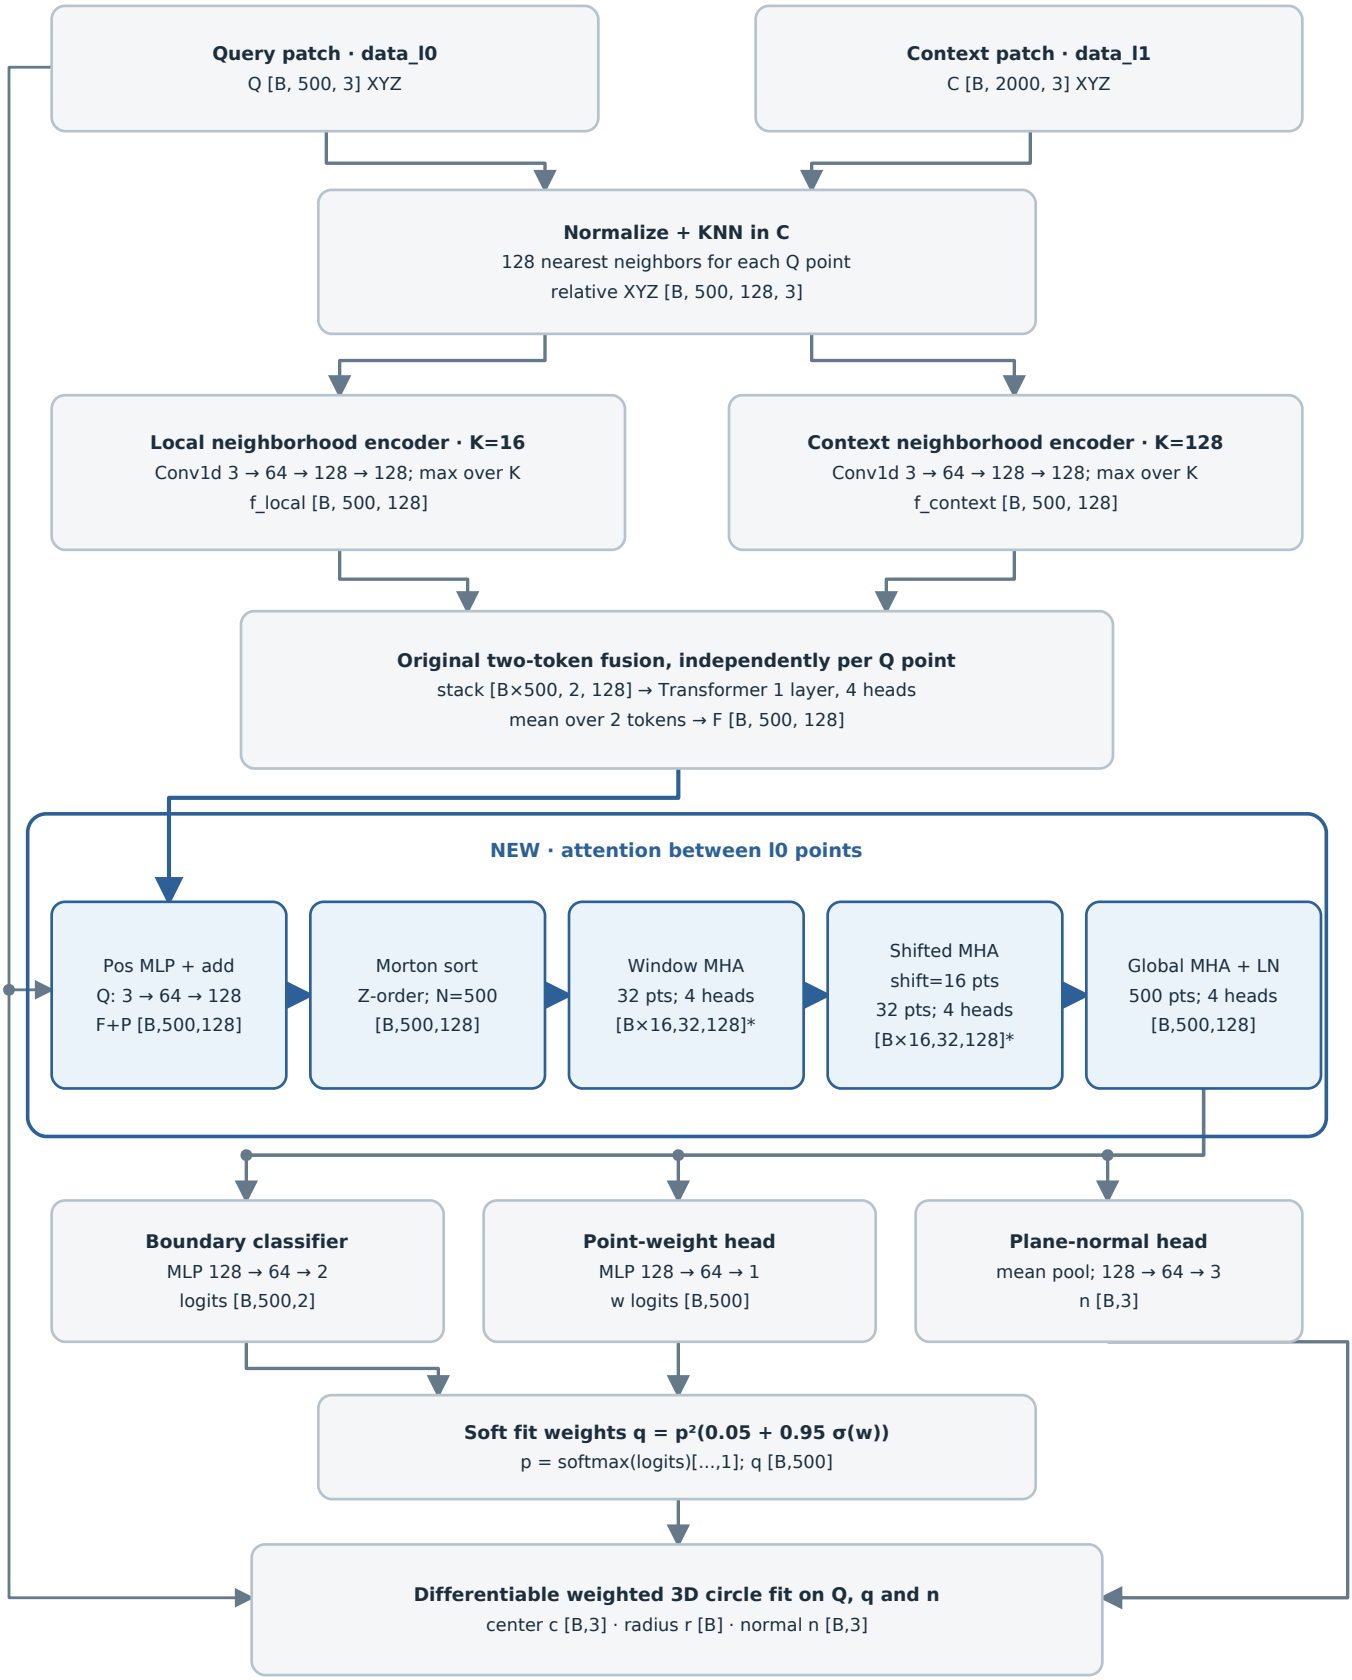

* 500 points are padded to 512 for 16 windows; 12 padding tokens are masked.

Each blue MHA block: LN → MHA → residual → LN → MLP(128→256→128) → residual.